


## AES Timing Attacks

Hardware and Software Design for Cryptographic Applications

April 11, 2013


## Ciphers as a Black Box

In theory, encryption (and decryption) implementations operate as black boxes.



## Information Leakage

In reality, it's hard to prevent additional information from being leaked at runtime.



## Side Channel Attacks

**Definition:** Any attack on a cryptosystem using information leaked given off as a byproduct of the physical implementation of the cryptosystem, rather than a theoretical weakness [1], is a *side channel attack*.

We focus on **timing attacks** for **software implementations** of AES.

We assume the attacker can *easily* capture this timing information.


# History of Timing Attacks on Cryptographic Primitives

## ■ LIST OF PAPERS FROM BONNEAU HISTORY

## Timing Attacks on AES

- Rijndael was deemed not susceptible to timing attacks in the AES contest
- AES targeted attacks can be based on *statistical* evidence [2].
  - Observation: The entire encryption time can be affected by the input bytes  $p_i^0 \oplus k_i^0$ 
    - Why?
      - Step 1: Capture timing data on a reference and target machine for each value of a particular input byte  $p_i^0 \oplus k_i^0 = x_i^0$
      - Step 2: Perform correlation between reference and target data
      - Step 3: Dance.
- Or they can be more targeted:
  - Exploit relationships between secret information of the primitive and known data.
  - This is the approach of Bonneau et al., among others.

# Cache Memory



$\langle x_i \rangle = \langle x_j \rangle$  are the higher bits of the data entry.

Data is pulled into cache based on the most-significant bits in its address.

## Cache Collisions Reveal Weaknesses

Let  $T_E(K, P)$  be the encryption time for a plaintext  $P$  using key  $K$ . Let  $\bar{T}_E(K) = \frac{1}{n} \sum_{i=1}^n T_E(K, P_i)$ , where  $P_i$  is a random plaintext from  $\{0|1\}^{128}$ .

**Cache-Collision Assumption [?].** For any pair of table lookups  $i, j$ , given a large enough number of *random* AES encryption that use the *same* key,  $\bar{T}_E(K)$  will be lower when  $\langle l_i \rangle = \langle l_j \rangle$  than when  $\langle l_i \rangle \neq \langle l_j \rangle$

Note: The table lookup indices must be *independent* for random plaintexts.



## Cache Collisions (cont'd)

Let  $a$  and  $b$  be two memory addresses looked up in memory. Let  $\langle a \rangle$  and  $\langle b \rangle$  denote the MSBs of  $a$  and  $b$ , respectively.

- Cache memory is organized into *lines*
  - MSB is mapped to the cache line index, LSB is mapped to the line (block) offset
- Reads on  $a$  and  $b$  cause a collision if  $\langle a \rangle = \langle b \rangle$  (assuming other memory reads have not evicted (or invalidated)  $a$  or  $b$  from the cache).
- If  $\langle a \rangle \neq \langle b \rangle$  then a cache collision *might* occur.
- We cannot say for certain whether or not the lower LSBs are equivalent...

## Attacks from Cache Collisions

That's it! We may now build an attack based on this result.

## LUT-Based Implementations

Let  $X^i$  be the state of AES at round  $i$ . With the exception of  $i = 10$ , we have:

$$\begin{aligned} X^{i+1} = & \{ T_0[x_0^i] \oplus T_1[x_5^i] \oplus T_2[x_{10}^i] \oplus T_3[x_{15}^i] \oplus \{k_0^i, k_1^i, k_2^i, k_3^i\}, \\ & T_0[x_4^i] \oplus T_1[x_9^i] \oplus T_2[x_{14}^i] \oplus T_3[x_3^i] \oplus \{k_4^i, k_5^i, k_6^i, k_7^i\}, \\ & T_0[x_8^i] \oplus T_1[x_{13}^i] \oplus T_2[x_2^i] \oplus T_3[x_7^i] \oplus \{k_8^i, k_9^i, k_{10}^i, k_{11}^i\}, \\ & T_0[x_{12}^i] \oplus T_1[x_1^i] \oplus T_2[x_6^i] \oplus T_3[x_{11}^i] \oplus \{k_{12}^i, k_{13}^i, k_{14}^i, k_{15}^i\} \} \end{aligned}$$

## First Round

- First round:  $x_i^0 = p_i \oplus k_i$
- With the T-box implementation,  $x_0^0$ ,  $x_4^0$ ,  $x_8^0$ , and  $x_{12}^0$  are used as indices into  $T_0$
- If we are looking for cache collisions, we must consider input bytes of the same T-box.

$$\begin{aligned}\langle x_i^0 \rangle = \langle x_j^0 \rangle &\Rightarrow \langle p_i \rangle \oplus \langle k_i \rangle = \langle p_j \rangle \oplus \langle k_j \rangle \\ &\Rightarrow \langle p_i \rangle \oplus \langle p_j \rangle = \langle k_i \rangle \oplus \langle k_j \rangle\end{aligned}$$

## First Round Attack Algorithm

---

### ALGORITHM 1: FirstRoundAttack( $N_s$ )

```
1:  $n \leftarrow 2^8 - 1$ 
2:  $T \leftarrow \text{array}[0 \dots n, 1 \dots n, 0 \dots n]$ 
3: for  $i = 0 \rightarrow N_s$  do
4:    $P \leftarrow \text{RandomBytes}(16)$ 
5:    $\text{start} \leftarrow \text{time}()$ 
6:    $C \leftarrow E_K(P)$ 
7:    $\text{end} \leftarrow \text{time}()$ 
8:    $tt \leftarrow (\text{start} - \text{end})$ 
9:    $T[i, j, \langle p_i \rangle \oplus \langle p_j \rangle] \leftarrow T[i, j, \langle p_i \rangle \oplus \langle p_j \rangle] + tt$ 
10: end for
11:  $T[i, j, \langle p_i \rangle \oplus \langle p_j \rangle] \leftarrow T[i, j, \langle p_i \rangle \oplus \langle p_j \rangle] / N_s$ 
12:  $mi, mj \leftarrow \min(t)$ 
13:  $\langle k_{mi} \rangle \oplus \langle k_{mj} \rangle \leftarrow \langle p_{mi} \rangle \oplus \langle p_{mj} \rangle$ 
```

▷ Use t-test if necessary

## First Round Attack Algorithm (cont'd)

$$\begin{aligned} X^{i+1} = & \{ T_0[x_0^i] \oplus T_1[x_5^i] \oplus T_2[x_{10}^i] \oplus T_3[x_{15}^i] \oplus \{k_0^i, k_1^i, k_2^i, k_3^i\}, \\ & T_0[x_4^i] \oplus T_1[x_9^i] \oplus T_2[x_{14}^i] \oplus T_3[x_3^i] \oplus \{k_4^i, k_5^i, k_6^i, k_7^i\}, \\ & T_0[x_8^i] \oplus T_1[x_{13}^i] \oplus T_2[x_2^i] \oplus T_3[x_7^i] \oplus \{k_8^i, k_9^i, k_{10}^i, k_{11}^i\}, \\ & T_0[x_{12}^i] \oplus T_1[x_1^i] \oplus T_2[x_6^i] \oplus T_3[x_{11}^i] \oplus \{k_{12}^i, k_{13}^i, k_{14}^i, k_{15}^i\} \} \end{aligned}$$

## First Round Attack Algorithm (cont'd)

$$\begin{aligned} X^{i+1} = & \{ \mathcal{T}_0[x_0^i] \oplus \mathcal{T}_1[x_5^i] \oplus \mathcal{T}_2[x_{10}^i] \oplus \mathcal{T}_3[x_{15}^i] \oplus \{k_0^i, k_1^i, k_2^i, k_3^i\}, \\ & \mathcal{T}_0[x_4^i] \oplus \mathcal{T}_1[x_9^i] \oplus \mathcal{T}_2[x_{14}^i] \oplus \mathcal{T}_3[x_3^i] \oplus \{k_4^i, k_5^i, k_6^i, k_7^i\}, \\ & \mathcal{T}_0[x_8^i] \oplus \mathcal{T}_1[x_{13}^i] \oplus \mathcal{T}_2[x_2^i] \oplus \mathcal{T}_3[x_7^i] \oplus \{k_8^i, k_9^i, k_{10}^i, k_{11}^i\}, \\ & \mathcal{T}_0[x_{12}^i] \oplus \mathcal{T}_1[x_1^i] \oplus \mathcal{T}_2[x_6^i] \oplus \mathcal{T}_3[x_{11}^i] \oplus \{k_{12}^i, k_{13}^i, k_{14}^i, k_{15}^i\} \} \end{aligned}$$

## First Round Limitation

We only know that  $\langle k_0 \rangle \oplus \langle k_4 \rangle = \Delta_1$ ,  $\langle k_0 \rangle \oplus \langle k_8 \rangle = \Delta_2$ ,  $\langle k_0 \rangle \oplus \langle k_{12} \rangle = \Delta_3$ ,  
 $\langle k_4 \rangle \oplus \langle k_8 \rangle = \Delta_4$ ,  $\langle k_4 \rangle \oplus \langle k_{12} \rangle = \Delta_5$ , and  $\langle k_8 \rangle \oplus \langle k_{12} \rangle = \Delta_6$ .

Of course, there exists **18** other equations we can derive for  $T_1$ ,  $T_2$ , and  $T_3$ .

We cannot determine the lower  $\log_2$  bits of each key... What to do now?



## The Last Round

When  $i = 10$ , the lookup table is just the S-box  $S$ . At this point, the ciphertext  $C$  is:

$$C = \{S[x_0^{10}] \oplus k_0^{10}, S[x_5^{10}] \oplus k_1^{10}, S[x_{10}^{10}] \oplus k_2^{10}, S[x_{15}^{10}] \oplus k_3^{10}, \\ S[x_4^{10}] \oplus k_5^{10}, S[x_9^{10}] \oplus k_6^{10}, S[x_{14}^{10}] \oplus k_7^{10}, S[x_3^{10}] \oplus k_7^{10}, \\ S[x_8^{10}] \oplus k_8^{10}, S[x_{13}^{10}] \oplus k_9^{10}, S[x_2^{10}] \oplus k_{10}^{10}, S[x_7^{10}] \oplus k_{11}^{10}, \\ S[x_{12}^{10}] \oplus k_{12}^{10}, S[x_1^{10}] \oplus k_{13}^{10}, S[x_6^{10}] \oplus k_{14}^{10}, S[x_{11}^{10}] \oplus k_{15}^{10}\}$$

## Final Round Collisions

Let  $x_s$  and  $x_t$  be two random bytes in the last round.

We will always have that  $c_i = k_i^{10} \oplus S[x_s]$  and  $c_j = k_j^{10} \oplus S[x_t]$ . If  $x_s = x_t$ , then a collision *will usually* occur and  $c_i = k_i^{10} \oplus \alpha$  and  $c_j = k_j^{10} \oplus \alpha$ .

Therefore,  $c_i \oplus c_j = k_i^{10} \oplus k_j^{10}$

## Final Round Misses

What if  $x_s \neq x_t$ ?

$c_i \oplus c_j \neq k_i^{10} \oplus k_j^{10} \Rightarrow$  two different values came out of the LUT!

The *S-box nonlinearity* means that the difference between  $S[x_s]$  and  $S[x_t]$  *does not* imply a similar difference between  $x_s$  and  $x_t$ .

If  $c_i \oplus c_j \neq k_i^{10} \oplus k_j^{10}$ , then  $x_s$  and  $x_t$  are two *random* values.

## Final Round Attack Algorithm

---

### ALGORITHM 2: LastRoundAttack( $N_s$ )

```

1:  $diffs \leftarrow [0, \dots, 2^n - 1, 0, \dots, 2^n - 1]$ 
2:  $T \leftarrow \text{array}[0 \dots 2^n - 1, 0 \dots 2^n - 1, 0 \dots 2^n - 1]$ 
3: for  $i = 0 \rightarrow N_s$  do
4:    $P \leftarrow \text{RandomBytes}(16)$ 
5:    $start \leftarrow \text{time}()$ 
6:    $C \leftarrow E_K(P)$ 
7:    $end \leftarrow \text{time}()$ 
8:    $tt \leftarrow (start - end)$ 
9:   for  $i, j$  in  $C$  do
10:     $T[i, j, C_i \oplus C_j] \leftarrow T[i, j, C_i \oplus C_j] + tt$ 
11:   end for
12: end for
13:  $T[i, j, \Delta_{i,j}] \leftarrow T[i, j, \Delta_{i,j}] / N_s$ 
14:  $\Delta'_{i,j} \leftarrow \min(T, i, j)$ 

```

▷ Use t-test if necessary **return** all  $\Delta'_{i,j}$

---

## Final Round Attack Algorithm (cont'd)

With knowledge of  $\Delta_{i,j}$  for all  $i,j$  such that  $\Delta_{i,j} = k_i^{10} \oplus k_j^{10}$  the attacker can now make informed guesses at the key

Thanks to Rijndael's invertible key schedule,  $[k^{10}]$  can be reverted back to  $[k^0]$ .

Done.

# Timing Attack Countermeasures

■ Keep?

## References

- 1 Joseph Bonneau and Ilya Mironov. Cache-collision timing attacks against AES. *Cryptographic Hardware and Embedded Systems CHES 2006*, Springer Berlin Heidelberg, (2006), 201-215.
- 2 Daniel J. Bernstein. Cache-timing attacks on AES. April 2005.  
<http://cr.yp.to/antiforgery/cachetiming-20050414.pdf>.